



FACT SHEET

DATA MODEL AND SCHEMA

- MongoDB is a **document-based data model**
- Data is stored in BSON (Binary JSON), which allows for a *schema-less design*
- Documents consist of key-value pairs
- Supports complex data structures like arrays and nested documents
- Unlike relational databases, collections do not enforce a fixed schema

CONSISTENCY AND REPLICATION

Theorem & Consistency

MongoDB is a NoSQL database and follows CAP theorem principles:

Partition Tolerance (P):

- Built to work in distributed environments

Consistency vs. Availability Trade-off:

- Strong consistency when reading from the primary node
- Eventual consistency when reading from secondary nodes

Tunable Consistency:

- Developers can choose read and write concerns to balance consistency and availability
- "Majority" read concern ensures data consistency across multiple nodes

SECURITY

- **Authentication & Authorization:**

- Supports username/password authentication
- Integrates with LDAP and x.509 certificate-based authentication
- Role-Based Access Control (RBAC) for fine-grained permissions

- **Encryption:**

- Encryption at Rest: Protects stored data
- Encryption in Transit: Uses TLS/SSL to secure client-server communication

- **Auditing & Compliance:** Provides detailed logging for security monitoring

SPECIFIC USE CASES

• **Rapid Development & Schema Flexibility:**

Ideal for startups and evolving applications

• **Big Data & Real-Time Analytics:**

Supports high-velocity data processing

• **Content Management Systems:**

Stores complex and unstructured data

• **IoT & Mobile Applications:**

Handles large-scale, high-speed data ingestion

• **Geospatial & Search Applications:**

Provides built-in geospatial indexing and full-text search

EXTRA FEATURES

- **Indexing:** Supports various types of indexes (compound, geospatial, text, etc.) to optimize queries
- **Aggregation Framework:** Provides a pipeline-based system for data transformation and analysis
- **Flexible Query Language:** Supports rich queries, including filtering, sorting, and regex searches
- **Change Streams:** Enables real-time data updates for applications
- **Replication and High Availability:** Uses replica sets for failover and data redundancy



CLUSTERS

Replica Set

- A replica set is a group of MongoDB servers that store identical copies of data.
- Purpose:
 - *High availability* (data is always accessible).
 - *Redundancy* (backup in case a server fails).
 - Handles failovers and maintenance with minimal downtime.
- Can handle read operations.

Sharded Cluster

- Also known as horizontal scaling.
- Data is split and distributed across multiple servers (shards).
- Purpose:
 - Scales read and write operations.
 - Useful when dealing with large datasets or high traffic.

OPLOG AND ELECTION PROCESS

Oplog:

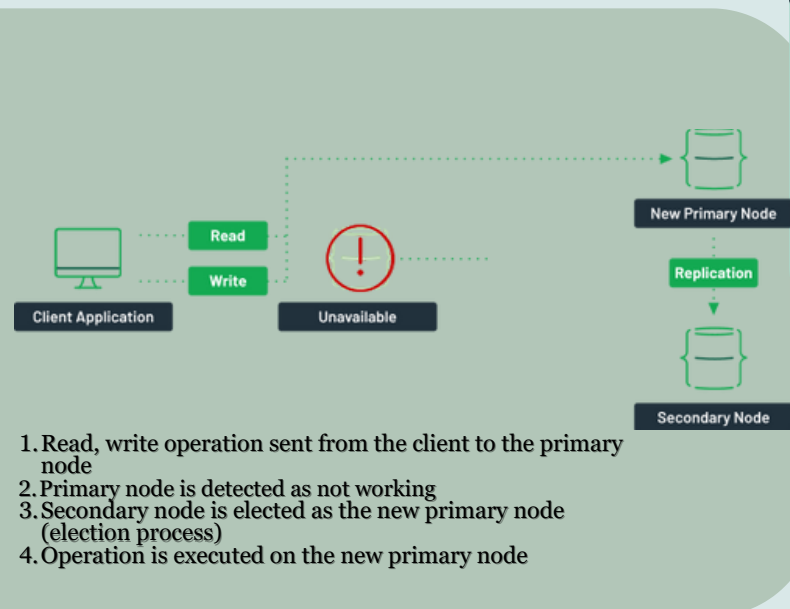
- special capped collection that stores a rolling log of operations.
- Used to sync secondary nodes with the primary.
- Dynamically resizes to avoid deleting important commits prematurely.

Election Process:

- Triggered when the primary becomes unavailable.
- A new primary is elected from secondary nodes.

Triggers include:

- Node failure
- Timeout > 10 seconds (default)
- Adding new nodes
- Initial setup or planned maintenance



HISTORY

In **2009**, MongoDB was officially released as an open-source project, allowing developers worldwide to leverage its features for free. The database quickly gained popularity due to its document-oriented design, scalability, and flexibility. Unlike traditional SQL databases, MongoDB provided a schema-less structure, making it particularly well-suited for rapidly evolving applications and big data workloads.

Economic information

MongoDB reported \$2.01 billion in revenue for fiscal year 2025, with its cloud service Atlas making up 68% of that, showing strong growth in the cloud database market. It holds a leading position among NoSQL databases and is used by major companies like Adobe, eBay, and Coinbase.

